


Verification Gap Engine

Three tracks now, not two. **Track A** only ever produces **Verified** or **Hallucinated** — a claim actually checked against a real source. **Track B** only ever produces **Unverifiable** — a claim that cited something, but the trail ran out. **Track C** is the one that closes the real gap: a claim that never cited anything at all — stated as flat fact, with no footnote to even test. It can only ever produce **Unsupported Assertion**. None of the three ever blend into each other.

v4 — adds Track C: everything up to now only ever looked at text that already carried a citation marker. That's a real blind spot — an AI can state something as fact without even attempting a footnote, which is easier to produce than a fake citation and had nothing in this pipeline to catch it. Track C branches directly off Stage 1, before the citation graph is even built, and scans all uncited body text for factual-sounding claims. It gets its own 3-stage pipeline, its own color, and its own bucket in the Scorecard — never merged into Unverifiable, because "cited but unbacked" and "never cited at all" are different failure modes with different fixes.



What Track C fixes: every stage before this version only ever looked at text that already had a footnote marker attached — which means a sentence stating something as flat fact, with no citation attempt at all, was invisible to the whole pipeline. It never reached the fork, because the fork only exists for citation rows, and this text was never a citation row. Track C branches directly off Stage 1's raw output, before the citation graph is even built, filters that uncited text for real factual claims (C1), scores how much it matters (C2), and labels it (C3) — a distinct fifth bucket, **Unsupported Assertion**, that can never be confused with Unverifiable. Unverifiable means "cited, but the trail ran out." Unsupported Assertion means "never even tried to cite anything." Different failure, different fix.