

Verification Gap Screens

Eight screens now, not seven — the gap review found real holes: no processing state between upload and results, only one verdict type ever demonstrated, and no view of what a human reviewer actually sees mid-decision. All fixed below.

v2 — fills every gap from the coverage review: added a **Processing** screen showing Stages 1–4 actually running. Citation Detail now shows all four verdict states, not just Verified — including a worked **Hallucinated** example (real audit had zero, so this one's marked illustrative), a **Pending** example with the reviewer confirm/override buttons that were missing before, and a **Track C** layout that's genuinely different — no source panel, because there's no source to show. Review Queue gained its own decision log, separate from Governance's policy log. Citation Detail examples now show a cache timestamp, surfacing the Source Cache store that had no UI presence at all.

1 Upload

Where an audit starts. The report-type dropdown matters more than it looks — it's what selects the Governance screen's threshold profile later.

verification-gap.app/new

1 · Upload

3 · Scorecard

4 · Citation detail

5 · Review queue

6 · Governance

7 · Export

8 · Monitoring

NEW AUDIT

Drop a PDF, DOCX, or HTML file here

Browse files

REPORT TYPE

Financial Services — Regulatory Outlook ▾

Start audit ▾

RECENT AUDITS

Regulatory Outlook for Financial Markets 202415 Aug 2026NEEDS REVIEW

— no other audits yet —

Why this screen matters: the report-type field is the only decision a user makes before the pipeline runs — everything after is automatic until Governance's thresholds gate the outcome.

2 Processing

New — closes the gap where Stages 1–4 had no visible state at all. Not a sidebar destination — a transient run a user watches, not navigates back to.

verification-gap.app/audits/reg-outlook-2024/run

1 · Upload

3 · Scorecard

4 · Citation detail

5 · Review queue

6 · Governance

7 · Export

8 · Monitoring

Analyzing document...

Regulatory Outlook for Financial Markets 2024 · 56 pages

✓ 1 · Extraction100 citation markers found · 49 bibliography entries

✓ 2 · Citation Graph Builder100 rows built · 49 matched to a source

✓ Fork49 → Track A · 51 → Track B · uncited text → Track C

⚙️ 3 · Claim Classifier — runningdispatching specialists to 49 Track A citations..

... Verdict Resolverpending

... Scorecardpending

Typically 2–4 minutes for a 56-page report

Why this screen matters: fetching and checking up to 49 real sources takes real time — a user watching a blank screen for minutes with no feedback is a worse product than one that's slightly slower but shows its work.

3 Scorecard

The headline view. Five separate counts, never blended into one score — and a gate status that's visible before anyone reads a single citation.

verification-gap.app/audits/reg-outlook-2024

1 · Upload

3 · Scorecard

4 · Citation detail

5 · Review queue

6 · Governance

7 · Export

8 · Monitoring

Regulatory Outlook for Financial Markets 202456 pages · 100 citations · audited 15 Aug 2026

NEEDS REVIEWExport ▾

49VERIFIED

0HALLUCINATED

51UNVERIFIABLE

0PENDING

—UNSUPPORTED

CITATIONS

#	PG	CLAIM	STATUS	MATERIALITY	
1	4	Klaas Knot, "inflation to stay above target..."	VERIFIED	Medium	View ▾
65	32	"£740bn (55%) of UK business lending via NBFIs"	UNVERIFIABLE	High	View ▾
81	44	"49% of portfolio managers... no vulnerable consumers"	UNVERIFIABLE	High	View ▾
96	49	Private credit "\$60bn in 2002 ... \$1.3T in H1 2022"	UNVERIFIABLE	High	View ▾
100	50	Deloitte's own "ESG in Private Capital Survey 2023"	UNVERIFIABLE	High	View ▾

Why this screen matters: the stacked bar makes the 51% gap visible in one glance, before anyone reads a single row — the whole point of the audit, surfaced above the fold.

4 Citation detail

Claim and source side by side, exactly which specialists ran, and what each found.

All four verdict states shown below — three of them are new.

STATE A — VERIFIED (REAL, FROM THE AUDIT)

verification-gap.app/audits/reg-outlook-2024/citations/1

1 · Upload

3 · Scorecard

4 · Citation detail

5 · Review queue

6 · Governance

7 · Export

8 · Monitoring

-- Scorecard

Citation #1 — Page 4VERIFIED

REPORT SAYS"Speech by Klaas Knot, President of the Netherlands Bank, June 2023..."

SOURCE SAYSDNB — "To improve is to change," Klaas Knot, 1 June 2023. Matches quoted content.

SPECIALISTS RUN1 · 5 OF 12 SHOWN — VIEW FULL ROSTER

✓ Existence ✓ Attribution ✓ Content-match ✓ Self-citation — Statistics (not triggered)

MATERIALITY: Medium · SOURCE CACHED: 15 Aug 2026, 09:14 — re-checked automatically by Monitoring

STATE B — HALLUCINATED (ILLUSTRATIVE — REAL AUDIT FOUND ZERO)

verification-gap.app/audits/example/citations/47

1 · Upload

3 · Scorecard

4 · Citation detail

5 · Review queue

6 · Governance

7 · Export

8 · Monitoring

-- Scorecard

Citation #47 — Page 19FABRICATED

REPORT SAYS"A 2023 survey by the Cambridge Centre for Financial Ethics found that 78% of retail investors distrust AI-driven advice."

SOURCE SAYS✗ No organization by this name found in any academic, regulatory, or press registry. Citation does not resolve.

SPECIALISTS RUN1 · 5 OF 12 SHOWN — VIEW FULL ROSTER

✗ Existence — no source resolves ✗ Attribution — organization not real — Content-match (skipped — nothing to compare)

✗ Statistics — 78% figure unverifiable ✓ Self-citation — n/a

MATERIALITY: High — quoted directly in Executive Summary

ESCALATED TO REVIEW QUEUE

Confirm fabricatedOverturn — mark verifiedAdd note

STATE C — PENDING, UNDER REVIEW (ILLUSTRATIVE — REAL AUDIT FOUND ZERO)

verification-gap.app/audits/example/citations/12

1 · Upload

3 · Scorecard

4 · Citation detail

5 · Review queue

6 · Governance

7 · Export

8 · Monitoring

-- Review queue

Citation #12 — Page 9PENDING

REPORT SAYS"Consumer confidence has fallen sharply across the sector this year."

SOURCE SAYSSource discusses a general downward trend but doesn't state "sharply," and covers a different 12-month window than the one implied.

VERDICT RESOLVER OUTPUTSuspicious — real source, partial support. Routed here automatically, not auto-resolved.

Assigned: B. Mangla · In review since 15 Aug 2026, 10:40

REVIEWER DECISION

Confirm verifiedConfirm fabricatedRequest more evidence

STATE D — UNSUPPORTED ASSERTION, TRACK C (ILLUSTRATIVE — TRACK C NOT YET RUN ON THE REAL REPORT)

verification-gap.app/audits/example/assertions/8

1 · Upload

3 · Scorecard

4 · Citation detail

5 · Review queue

6 · Governance

7 · Export

8 · Monitoring

-- Scorecard

Assertion #8 — Page 21 (no citation number — none was ever attached)UNSUPPORTED

CLAIM (UNCITED)"Regulators are increasingly concerned about this trend."

No "Source says" panel here — there's nothing to show. No citation was ever attempted for this claim.

SPECIALISTS RUN

✓ C1 · Claim-vs-narration filter — real factual claim, not narration ✓ C2 · Materiality tag

MATERIALITY: Low — supporting sentence, not a headline claim

Request citation be addedAccept as house style

Full specialist roster — all 12 (click to collapse)

CODE	SPECIALIST	TRIGGERS ON	CHECKS
A1	Existence	Always	Does the source resolve at all
A2	Attribution	Always	Right org · author · date
A3	Self-citation	Always	Cites its own domain?
A4	Content-match	Always	Does the source actually say this claim (bought: Patronus/Vectara)
A5	Author-verify	Named person + quote	Was this quote really said, right date/venue
A6	Statistics	\$ or % figure	Does the number match the source
A7	Chain-of-custody	Secondary/news source	Primary source, or laundered through one
A8	Framing-fidelity	Caveat in the source	Was the caveat dropped when the report used the fact
A9	Consistency	Repeated fact in-doc	Do the two instances agree with each other
A10	Derived-calc	Report's own math	Is the calculation shown and correct
A11	Chart-data	Chart/graph caption	Does the visualized data match the source
A12	Legal-citation	Law/regulation reference	Case or statute verified against a legal registry

A typical citation triggers 2–5 of these, not all 12 — that's the router working as intended, not a coverage gap. Track C runs a separate, 2-item roster (Claim-vs-narration filter, Materiality tag), since there's no source to run the other 12 against.

Why this screen matters now: a reviewer can finally see what it actually looks like when the system catches something — not just when it confirms something's fine.

5 Review queue

Escalated items only — the human-run stage. Now includes its own decision log, separate from Governance's policy log.

verification-gap.app/audits/reg-outlook-2024/review

1 · Upload

3 · Scorecard

4 · Citation detail

5 · Review queue

6 · Governance

7 · Export

8 · Monitoring

All (3)Suspicious (0)Pending (0)High-materiality Unverifiable (3)

☐	#	STATUS	CLAIM	ASSIGNED	ACTION
☐	65	UNVERIFIABLE	"£740bn (55%) of UK business lending via NBFIs"	Unassigned	Review
☐	81	UNVERIFIABLE	"49% of portfolio managers... no vulnerable consumers"	Unassigned	Review
☐	96	UNVERIFIABLE	Private credit "\$60bn in 2002 ... \$1.3T in H1 2022"	Unassigned	Review

Assign to meMark reviewed

DECISION LOG (ILLUSTRATIVE)

14 Aug 2026 — B. Mangla — Citation #47 (Suspicious) — confirmed Fabricated, escalated to comms

13 Aug 2026 — B. Mangla — Citation #12 (Suspicious) — confirmed Verified

Why this screen matters: only 3 of the 51 Unverifiable citations show up here — the ones tagged high-materiality by Stage 7. The decision log is what makes a reviewer's call auditable after the fact, not just in the moment.

6 Governance

Pure policy, no model — the coverage thresholds and zero-tolerance rules a firm sets, applied automatically to every future audit of that report type.

verification-gap.app/settings/governance

1 · Upload

3 · Scorecard

4 · Citation detail

5 · Review queue

6 · Governance

7 · Export

8 · Monitoring

COVERAGE THRESHOLDS — REGULATORY OUTLOOK PROFILE

Max % Unverifiable before Hold40%

This report is at 51% — over threshold

Max % Unsupported before Hold15%

ZERO-TOLERANCE RULES

☒ No Hallucinated citations in Executive Summary☒ No Hallucinated citations behind a pull-quote

Save policy

POLICY CHANGE LOG

12 Aug 2026 — B. Mangla — set Unverifiable threshold to 40% (was 55%)

Why this screen matters: this is the one screen with no AI in it at all — just the configured rule that turns "51% unverifiable" into an automatic Hold, and an audit trail of who set that number and when. Its own log is policy changes, not reviewer decisions — that's Screen 5's job.

7 Export

What actually leaves the system — a shareable report with a redaction toggle, so compliance sees everything and an external reader sees a summary.

verification-gap.app/audits/reg-outlook-2024/export

1 · Upload

3 · Scorecard

4 · Citation detail

5 · Review queue

6 · Governance

7 · Export

8 · Monitoring

REPORT PREVIEW

Regulatory Audit — Regulatory Outlook for Financial Markets 202449 Verified · 0 Hallucinated · 51 Unverifiable · 0 Pending

EVIDENCE APPENDIX

#65 — Unverifiable — p.32#81 — Unverifiable — p.44#96 — Unverifiable — p.49... 48 more

EXPORT AS

PDFExcelJSON (API)

VISIBILITY

● Full evidence (internal)○ Summary only (external)

Why this screen matters: "Summary only" exists because not every recipient of this report should see the same depth of internal evidence — the redaction toggle is a real requirement, not an afterthought.

8 Monitoring

The library view — trends across every report a firm has run through the system, and what's changed since a citation was first checked.

verification-gap.app/monitoring

1 · Upload

3 · Scorecard

4 · Citation detail

5 · Review queue

6 · Governance

7 · Export

8 · Monitoring

HALLUCINATION RATE — LAST 6 REPORTS

SELF-CITATION RATE — LAST 6 REPORTS

ALERTS

▲ 2 sources have changed since last check — view

AUDITED REPORTS

REPORT	AUDITED	LAST RE-CHECK	STATUS
Regulatory Outlook for Financial Markets 2024	15 Aug 2026	15 Aug 2026	UP TO DATE

Why this screen matters: a rising self-citation trend line is invisible from any single-document audit — this is the one screen where the value only shows up after a firm has run more than one report through the system.

Every real number still comes from the Deloitte Regulatory Outlook 2024 audit. The three new example states (Fabricated, Pending, Unsupported) are marked illustrative because the real audit never triggered them — that's a fact about the report, not something smoothed over here.